



Concurrency Control

Instructor : Nitesh Kumar Jha

ITER,S'O'A(DEEMED TO BE
UNIVERSITY)

Concurrency Control

- It is a mechanism to ensure isolation in a concurrent execution scenario
- Achieved using the concept of **mutual exclusion**
 - i.e. while one transaction is accessing a data item, no other transaction is allowed to modify that data item.
- Mutual exclusion is achieved using logical **locks**.
 - Locks are granted/revoked by a central **concurrency control manager**.
 - i.e. Transactions request it to grant a lock

Locks

- Data items can be locked in two modes :
 - *Shared*: Data item can only be read. Requested using **lock-S** instruction. It can be shared with other transactions
 - *Exclusive*: Data item can be both read as well as written. It is requested using **lock-X** instruction. It can't be shared with other transactions

- Lock-compatibility matrix

- Shared locks may be granted to multiple transactions simultaneously.
- Exclusive locks can't be granted to multiple transactions simultaneously

	S	X
S	true	false
X	false	false

Locks

- Let A and B are two accounts that are accessed by transactions T_1 and T_2 . Transaction T_1 transfers \$50 from account B to account A. Transaction T_2 displays the total amount of money in accounts A and B i.e., sum $A+B$.
- If these transaction are executed serially, $T_1 T_2$ or $T_2 T_1$, It will be consistent.
- What happens if these are executed concurrently???

T_1	T_2
Lock-X(B)	Lock-S(A)
Read(B)	Read(A)
$B=B-50$	Unlock(A)
Write(B)	Lock-S(B)
Unlock(B)	Read(B)
Lock-X(A)	Unlock(B)
Read(A)	Display (A+B)
$A=A+50$	
Write(A)	
Unlock(A)	

Locking Example

T ₁	T ₂	CC Manager	T ₁	T ₂	CC Manager
Lock-X(B)	A=500	Grant-X(B, T ₁)		Read(A)	Grant-S(A, T ₂)
	B=1000			Unlock(A)	1450
Read(B)	1000			Display (A+B)	
B=B-50	950				
Write(B)		Grant-S(B, T ₂)	Lock-X(A)		Grant-X(A, T ₁)
Unlock(B)			Read(A)		
	Lock-S(B)		A=A+50		
	Read(B)		Write(A)		
	Unlock(B)		Unlock(B)		
	Lock-S(A)				

- Produces inconsistent result
- Transactions must hold the lock on a data item till it access that item.
- It is not necessarily desirable for a transaction to unlock a data item immediately after its final access of that item, since the serializability may not be ensured.

Delayed Unlocking

- Unlocking is delayed to the end of the transactions
- This schedule produces consistent result

T_3 : lock-X(B);
read(B);
 $B := B - 50$;
write(B);
lock-X(A);
read(A);
 $A := A + 50$;
write(A);
unlock(B);
unlock(A).

T_4 : lock-S(A);
read(A);
lock-S(B);
read(B);
display($A + B$);
unlock(A);
unlock(B).

- But this technique may lead to an undesirable scenario called deadlock

Deadlock due to Hold-and-Wait

T_5	T_6	CC Manager
Lock-X(B)		Grant-X(B, T_1)
Read(B)		
$B = B - 50$		
Write(B)	Lock-S(A)	Grant-S(A, T_2)
	Read(A)	
Lock-X(A)	Lock-S(B)	

Deadlock

- ❑ When a transaction (T_5) delays unlocking on its locked data items (B) and requests to acquire a lock on new data items (A) that is already locked by another transaction (T_6)
- ❑ This is called a Hold-and-Wait situation

Deadlock due to Hold-and-Wait

T_5	T_6	CC Manager
Lock-X(B)		Grant-X(B, T_1)
Read(B)		
B=B-50		
Write(B)		
	Lock-S(A)	
		Grant-S(A, T_2)
	Read(A)	
	Lock-S(B)	
Lock-X(A)		
Deadlock		

- Deadlock is a state where neither of these transactions can ever proceed with its normal execution. It can be resolved by forcibly rolling back one or more participating transactions.
- Lock based concurrency control needs the transaction to follow a set of rules called **locking protocol**
- A locking protocol is a set of rules followed by all transactions while requesting and releasing locks.
- Locking protocols restrict the set of possible schedules.

The Two-Phase Locking Protocol

- This protocol ensures conflict-serializable schedules.
- It requires the transactions execute in two phases
- Phase 1: Growing Phase
 - Transaction may obtain locks
 - Transaction may not release locks
- Phase 2: Shrinking Phase
 - Transaction may release locks
 - Transaction may not obtain locks
- This protocol assures serializability.

Lock Conversions

- Two-phase locking with lock conversions:
 - First Phase:
 - can acquire a lock-S on item
 - can acquire a lock-X on item
 - can convert a lock-S to a lock-X (upgrade)
 - Second Phase:
 - can release a lock-S
 - can release a lock-X
 - can convert a lock-X to a lock-S (downgrade)
- This protocol assures serializability. But still relies on the programmer to insert the various locking instructions.

Automatic Acquisition of Locks

Not IMP

- A transaction T_i issues the standard read/write instruction, without explicit locking calls.
- The operation **read**(D) is processed as:
 - if** T_i has a lock on D
 - then**
 - read(D)
 - else begin**
 - if necessary wait until no other transaction has a **lock-X** on D
 - grant T_i a **lock-S** on D ;
 - read(D)
 - end**

Automatic Acquisition of Locks (Cont.)

Not IMP

- **write**(D) is processed as:

if T_i has a **lock-X** on D

then

write(D)

else begin

if necessary wait until no other transaction has any lock on D ,

if T_i has a **lock-S** on D

then

upgrade lock on D to **lock-X**

else

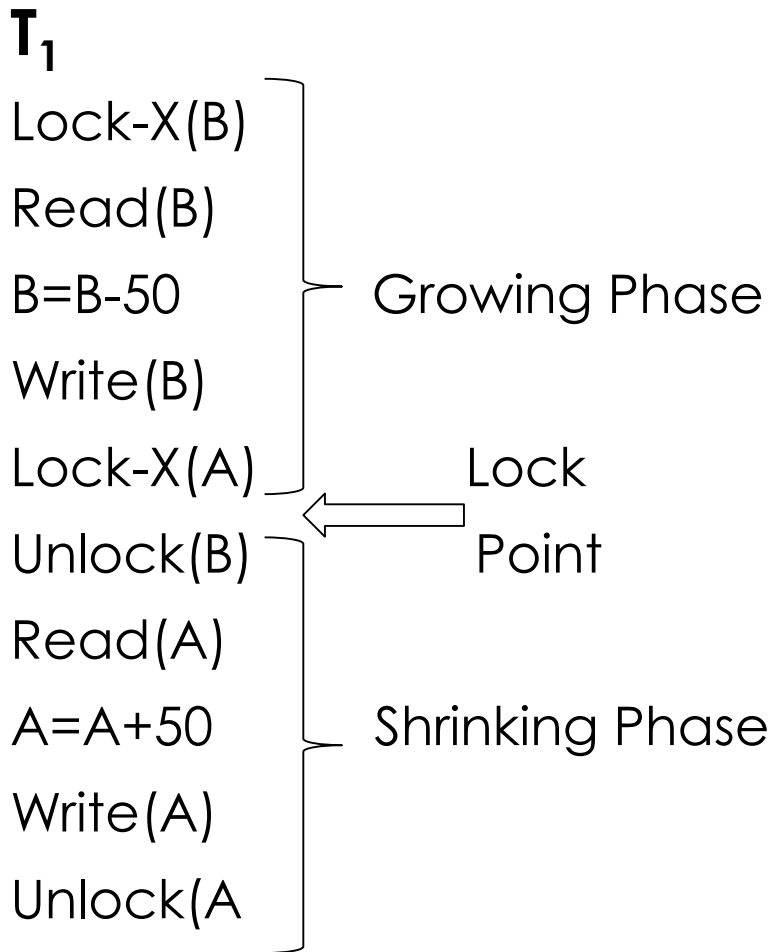
grant T_i a **lock-X** on D

write(D)

end;


- All locks are released after commit or abort

The Two-Phase Locking Protocol



- ❑ The point in the schedule where transaction has obtained its final lock(end of its growing phase) is called **lock point** of the transaction.
- ❑ Now transactions can be ordered according to their lock points.

Two Phase Locking Example

T ₁	T ₂	CC Manager	T ₁	T ₂	CC Manager
Lock-X(B)		Grant-X(B, T ₁)		Read(A)	
Read(B)	Not proceed further			Display (A+B)	1450
B=B-50	So T₁ Will proceed first then T₂			Unlock(B)	
Write(B)				Unlock(A)	
	Lock-S(B)	Grant-S(B, T ₂)	Lock-X(A)		Grant-X(A, T1)
	Read(B)		Read(A)		
	Lock-S(A)	Grant-S(A, T ₂)	A=A+50		
			Write(A)		
			Unlock(B)		
			Unlock(A)		

- Produces consistent result as transaction will be serial T₁ -> T₂
- Transactions holds the lock on a data item till the last.

Properties of 2PL

- 2PL ensures conflict serializability. The contributing transactions are isolated w.r.t. the lock point. (point at which growing phase ends and shrinking phase starts)
- It does not ensure deadlock free execution. In the event of deadlock participating transactions are rolled back. Consider

Previous Example here ----->

T ₅	T ₆	CC Manager
Lock-X(B)		Grant-X(B, T ₁)
Read(B)		
B=B-50		
Write(B)		
	Lock-S(A)	
	Read(A)	Grant-S(A, T ₂)
	Lock-S(B)	
Lock-X(A)		
Deadlock		

- It ensures recoverability but does not safeguard against cascading rollback.

Deadlocks

- Consider the partial schedule

T_3	T_4
lock-x (B)	
read (B)	
$B := B - 50$	
write (B)	
	lock-s (A)
	read (A)
	lock-s (B)
lock-x (A)	

- Neither T_3 nor T_4 can make progress — executing **lock-S(B)** causes T_4 to wait for T_3 to release its lock on B , while executing **lock-X(A)** causes T_3 to wait for T_4 to release its lock on A .
- Such a situation is called a **deadlock**.
 - To handle a deadlock one of T_3 or T_4 must be rolled back and its locks released.

Deadlocks (Cont.)

- The potential for deadlock exists in most locking protocols. Deadlocks are a necessary evil.
- When a deadlock occurs there is a possibility of cascading roll-backs.
- Cascading roll-back is possible under two-phase locking. To avoid this, we must follow a modified protocol called **strict two-phase locking** -- a transaction must hold all its exclusive locks till it commits/aborts.

Strict two Phase locking Protocol

- **Strict two phase locking** is an enhanced 2PL that ensures cascadeless recovery.
- Strict 2PL demands that not only the locking and unlocking be in two phases, all the exclusive mode locks must be held by the transaction till the transaction commits.
- This requirement ensured that any data written by uncommitted transactions are locked in Exclusive mode until the transaction commits, preventing any other transaction from reading the data.
- **Rigorous two-phase locking** is even stricter. Here, *all* locks are held till commit/abort. In this protocol transactions can be serialized in the order in which they commit.

Deadlocks (Cont.)

- Strict Two-phase locking *does not* ensure freedom from deadlocks.
- In addition to deadlocks, there is a possibility of **starvation**.
- **Starvation** occurs if the concurrency control manager is badly designed. For example:
 - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
 - The same transaction is repeatedly rolled back due to deadlocks.
- Concurrency control manager can be designed to prevent starvation.

Timestamp-Based Protocols

- Each transaction is issued a timestamp when it enters the system. If an old transaction T_i has time-stamp $TS(T_i)$, a new transaction T_j is assigned time-stamp $TS(T_j)$ such that $TS(T_i) < TS(T_j)$.
- The protocol manages concurrent execution such that the time-stamps determine the serializability order.
- In order to assure such behavior, the protocol maintains for each data Q two timestamp values:
 - **W-timestamp**(Q) is the largest time-stamp of any transaction that executed **write**(Q) successfully.
 - **R-timestamp**(Q) is the largest time-stamp of any transaction that executed **read**(Q) successfully.

Timestamp-Based Protocols (Cont.)


- The timestamp ordering protocol ensures that any conflicting **read** and **write** operations are executed in timestamp order.
- Suppose a transaction T_i issues a **read**(Q)
 1. If $TS(T_i) \leq \mathbf{W}\text{-timestamp}(Q)$, then T_i needs to read a value of Q that was already overwritten by a younger transaction.
 - Hence, the **read** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) \geq \mathbf{W}\text{-timestamp}(Q)$, then the **read** operation is executed, and $\mathbf{R}\text{-timestamp}(Q)$ is set to $\mathbf{max}(\mathbf{R}\text{-timestamp}(Q), TS(T_i))$.

Timestamp-Based Protocols (Cont.)

- Suppose that transaction T_i issues **write**(Q).
 1. If $TS(T_i) < R\text{-timestamp}(Q)$, then the value of Q that T_i is producing was needed previously, and the system assumed that that value would never be produced.
 - Hence, the **write** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of Q.
 - Hence, this **write** operation is rejected, and T_i is rolled back.
 3. Otherwise, the **write** operation is executed, and $W\text{-timestamp}(Q)$ is set to $TS(T_i)$.

Correctness of Timestamp-Ordering Protocol

- The timestamp-ordering protocol guarantees serializability since all the arcs in the precedence graph are of the form:



Thus, there will be no cycles in the precedence graph

- Timestamp protocol ensures freedom from deadlock as no transaction ever waits.
- But the schedule may not be cascade-free, and may not even be recoverable.

Recoverability and Cascade Freedom

- Problem with timestamp-ordering protocol:
 - Suppose T_i aborts, but T_j has read a data item written by T_i
 - Then T_j must abort; if T_j had been allowed to commit earlier, the schedule is not recoverable.
 - Further, any transaction that has read a data item written by T_j must abort
 - This can lead to cascading rollback --- that is, a chain of rollbacks
- Solution 1:
 - A transaction is structured such that its writes are all performed at the end of its processing
 - All writes of a transaction form an atomic action; no transaction may execute while a transaction is being written
 - A transaction that aborts is restarted with a new timestamp
- Solution 2: Limited form of locking: wait for data to be committed before reading it
- Solution 3: Use commit dependencies to ensure recoverability

End of Lecture
Thank You